

In [1]:

```
import os
import numpy as np
import pandas as pd
from PIL import Image

import lucid
import lucid.nn as nn
import lucid.nn.functional as F
import lucid.optim as optim
import lucid.transforms as transforms

from lucid.data import Dataset, DataLoader
from lucid.models.util import summarize

import matplotlib.pyplot as plt
from tqdm import tqdm
```

In [2]:

```
class CelebADataset(Dataset):
    def __init__(self, root_dir, attr_path="list_attr_celeba.txt", transform=None):
        self.root_dir = os.path.join(root_dir, "img_align_celeba")
        self.attr_path = os.path.join(root_dir, attr_path)
        self.transform = transform

        with open(self.attr_path, 'r') as f:
            lines = f.readlines()
            self.attr_names = lines[1].strip().split()
            self.attrs = pd.read_csv(self.attr_path, skiprows=2, delim_whitespace=True)
            self.attrs = self.attrs.astype({col: int for col in self.attrs.columns})
            self.attrs = (self.attrs + 1) // 2

            self.filenames = self.attrs.index.tolist()

    def __len__(self):
        return len(self.filenames)

    def __getitem__(self, idx):
        img_name = self.filenames[idx]
        img_path = os.path.join(self.root_dir, img_name)
        image = Image.open(img_path).convert("RGB")
        image = np.array(image)
        image = image.transpose((2, 0, 1))

        attr = self.attrs.loc[img_name].values.astype("float32")

        if self.transform:
            image = self.transform(image)

        return image, attr
```

In [3]:

```
transform = transforms.Compose([
    transforms.ToTensor(dtype=lucid.Int16),
    transforms.Resize(128),
    transforms.CenterCrop(128),
])

train_dataset = CelebADataset(root_dir=".data/celeba/", transform=transform)
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
```

```
/var/folders/bx/znht1lyx04j9q74ngjv7xvwh0000gn/T/ipykernel_34429/793879903.py:10: FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed in a future version. Use ``sep='\s+'`` instead
    self.attrs = pd.read_csv(self.attr_path, skiprows=2, delim_whitespace=True)
```

```
e, index_col=0)
```

```
In [4]: for images, attrs in train_loader:
        print(images.shape)
        print(attrs.shape)
        break
```

```
(64, 3, 128, 128)
(64, 40)
```

```
In [5]: class VAE(nn.Module):
        def __init__(self, latent_dim):
            super().__init__()
            hidden_dims = [32, 64, 128, 256, 512]
            self.final_dim = hidden_dims[-1]
            self.latent_dim = latent_dim
            in_channels = 3

            modules = []
            for h_dim in hidden_dims:
                modules.append(
                    nn.Sequential(
                        nn.Conv2d(in_channels, h_dim, kernel_size=3, stride=2,
                                nn.BatchNorm2d(h_dim),
                                nn.LeakyReLU(),
                                )
                    )
                in_channels = h_dim

            self.encoder = nn.Sequential(*modules)
            self.fc_mu = nn.Linear(hidden_dims[-1] * 4 * 4, self.latent_dim)
            self.fc_logvar = nn.Linear(hidden_dims[-1] * 4 * 4, self.latent_dim)

            modules = []
            self.decoder_input = nn.Linear(self.latent_dim, hidden_dims[-1] * 4 * 4)

            hidden_dims.reverse()
            for i in range(len(hidden_dims) - 1):
                modules.append(
                    nn.Sequential(
                        nn.ConvTranspose2d(
                            hidden_dims[i],
                            hidden_dims[i + 1],
                            kernel_size=3,
                            stride=2,
                            padding=1,
                            output_padding=1,
                        ),
                        nn.BatchNorm2d(hidden_dims[i + 1]),
                        nn.LeakyReLU(),
                    )
                )

            self.decoder = nn.Sequential(*modules)
            self.final_layer = nn.Sequential(
                nn.ConvTranspose2d(
                    hidden_dims[-1],
                    hidden_dims[-1],
                    kernel_size=3,
                    stride=2,
                    padding=1,
                    output_padding=1,
                ),
                nn.BatchNorm2d(hidden_dims[-1]),
                nn.LeakyReLU(),
                nn.Conv2d(hidden_dims[-1], 3, kernel_size=3, padding=1),
                nn.Sigmoid(),
            )

        def encode(self, x):
            ret = self.encoder(x)
```

```

ret = lucid.flatten(ret, axis=1)
mu = self.fc_mu(ret)
logvar = self.fc_logvar(ret).clip(-10.0, 10.0)
return mu, logvar

def reparameterize(self, mu, logvar):
    std = lucid.exp(0.5 * logvar)
    eps = lucid.random.randn(*std.shape, device=std.device)
    return eps * std + mu

def decode(self, z):
    ret = self.decoder_input(z)
    ret = ret.reshape(-1, self.final_dim, 4, 4)
    ret = self.decoder(ret)
    ret = self.final_layer(ret)
    return ret

def forward(self, x):
    mu, logvar = self.encode(x)
    z = self.reparameterize(mu, logvar)
    return self.decode(z), mu, logvar

```

In [6]:

```

device = "gpu"

model = VAE(latent_dim=128).to(device)
optimizer = optim.Adam(model.parameters(), lr=1e-3)

summarize(model, input_shape=(1, 3, 128, 128), test_backward=True)

```

Summary of VAE

Layer	Input Shape	Output Shape
Parameter Size		
=====		
VAE	(1, 3, 128, 128)	None
6,303,075		
└─ Sequential	(1, 32, 64, 64)	(1, 3, 128, 128)
10,179		
└─ Sigmoid	(1, 3, 128, 128)	(1, 3, 128, 128)
-		
└─ Conv2d	(1, 32, 128, 128)	(1, 3, 128, 128)
867		
└─ LeakyReLU	(1, 32, 128, 128)	(1, 32, 128, 128)
-		
└─ BatchNorm2d	(1, 32, 128, 128)	(1, 32, 128, 128)
64		
└─ ConvTranspose2d	(1, 32, 64, 64)	(1, 32, 128, 128)
9,248		
└─ Sequential	(1, 512, 4, 4)	(1, 32, 64, 64)
1,568,160		
└─ Sequential	(1, 64, 32, 32)	(1, 32, 64, 64)
18,528		
└─ LeakyReLU	(1, 32, 64, 64)	(1, 32, 64, 64)
-		
└─ BatchNorm2d	(1, 32, 64, 64)	(1, 32, 64, 64)
64		
└─ ConvTranspose2d	(1, 64, 32, 32)	(1, 32, 64, 64)
18,464		
└─ Sequential	(1, 128, 16, 16)	(1, 64, 32, 32)
73,920		
└─ LeakyReLU	(1, 64, 32, 32)	(1, 64, 32, 32)
-		

```

|   |--- BatchNorm2d                (1, 64, 32, 32)      (1, 64, 32, 32)
128
|   |--- ConvTranspose2d            (1, 128, 16, 16)    (1, 64, 32, 32)
73,792
|--- Sequential                    (1, 256, 8, 8)      (1, 128, 16, 16)
295,296
|   |--- LeakyReLU                 (1, 128, 16, 16)    (1, 128, 16, 16)
-
|   |--- BatchNorm2d                (1, 128, 16, 16)    (1, 128, 16, 16)
256
|   |--- ConvTranspose2d            (1, 256, 8, 8)      (1, 128, 16, 16)
295,040
|--- Sequential                    (1, 512, 4, 4)      (1, 256, 8, 8)
1,180,416
|   |--- LeakyReLU                 (1, 256, 8, 8)      (1, 256, 8, 8)
-
|   |--- BatchNorm2d                (1, 256, 8, 8)      (1, 256, 8, 8)
512
|   |--- ConvTranspose2d            (1, 512, 4, 4)      (1, 256, 8, 8)
1,179,904
|--- Linear                        (1, 128)             (1, 8192)
1,056,768
|--- Linear                        (1, 8192)            (1, 128)
1,048,704
|--- Linear                        (1, 8192)            (1, 128)
1,048,704
|--- Sequential                    (1, 3, 128, 128)    (1, 512, 4, 4)
1,570,560
|--- Sequential                    (1, 256, 8, 8)      (1, 512, 4, 4)
1,181,184
|   |--- LeakyReLU                 (1, 512, 4, 4)      (1, 512, 4, 4)
-
|   |--- BatchNorm2d                (1, 512, 4, 4)      (1, 512, 4, 4)
1,024
|   |--- Conv2d                    (1, 256, 8, 8)      (1, 512, 4, 4)
1,180,160
|--- Sequential                    (1, 128, 16, 16)    (1, 256, 8, 8)
295,680
|   |--- LeakyReLU                 (1, 256, 8, 8)      (1, 256, 8, 8)
-
|   |--- BatchNorm2d                (1, 256, 8, 8)      (1, 256, 8, 8)
512
|   |--- Conv2d                    (1, 128, 16, 16)    (1, 256, 8, 8)
295,168
|--- Sequential                    (1, 64, 32, 32)     (1, 128, 16, 16)
74,112
|   |--- LeakyReLU                 (1, 128, 16, 16)    (1, 128, 16, 16)
-
|   |--- BatchNorm2d                (1, 128, 16, 16)    (1, 128, 16, 16)
256
|   |--- Conv2d                    (1, 64, 32, 32)     (1, 128, 16, 16)
73,856
|--- Sequential                    (1, 32, 64, 64)     (1, 64, 32, 32)
18,624
|   |--- LeakyReLU                 (1, 64, 32, 32)     (1, 64, 32, 32)
-
|   |--- BatchNorm2d                (1, 64, 32, 32)     (1, 64, 32, 32)
128
|   |--- Conv2d                    (1, 32, 64, 64)     (1, 64, 32, 32)
18,496
|--- Sequential                    (1, 3, 128, 128)    (1, 32, 64, 64)
960
|   |--- LeakyReLU                 (1, 32, 64, 64)     (1, 32, 64, 64)
-
|   |--- BatchNorm2d                (1, 32, 64, 64)     (1, 32, 64, 64)
64
|   |--- Conv2d                    (1, 3, 128, 128)    (1, 32, 64, 64)
896

```

```

=====
Total Layers(Submodules): 47
Total Parameters: 6,303,075 (6.30M)

```

Total FLOPs: 19,214,496 (19.21M)

=====

```
In [7]: def normalize(img):
        ret = img.astype(lucid.Float32) / 255.0
        return ret
```

```
In [8]: def loss_function(recon_x, x, mu, logvar):
        mse = F.mse_loss(recon_x, x, reduction="mean")
        kld = -0.5 * lucid.mean(1 + logvar - mu ** 2 - lucid.exp(logvar))
        kld_weight = 0.0005
        loss = mse + kld_weight * kld
        return loss
```

```
In [9]: def train(num_epochs):
        losses = []
        model.train()
        for epoch in range(num_epochs):
            batch_losses = []
            pbar = tqdm(train_loader, desc=f"Epoch {epoch + 1}/{num_epochs}")

            for data, _ in pbar:
                data = data.to(device)
                data = normalize(data)
                optimizer.zero_grad()
                recon_data, mu, logvar = model(data)

                loss = loss_function(recon_data, data, mu, logvar)
                loss.eval()
                loss.backward()
                optimizer.step()

                loss_value = loss.item()
                batch_losses.append(loss_value)

            avg_loss = sum(batch_losses) / len(batch_losses)
            pbar.set_postfix(loss=loss_value, avg_loss=avg_loss)

            losses.extend(batch_losses)

        return losses
```

```
In [ ]: losses = train(num_epochs=10)
```

```
In [ ]: plt.plot(losses, lw=1, label="Recon + KL-Div")
        plt.xlabel("Iteration")
        plt.ylabel("Loss")
        plt.title("VAE Training for CelebA Dataset")
        plt.grid(alpha=0.3)
        plt.legend()
        plt.tight_layout()
```

```
Out[ ]: '\nplt.plot(losses, lw=1, label="Recon + KL-Div")\nplt.xlabel("Iteration")\n
        plt.ylabel("Loss")\nplt.title("VAE Training for CelebA Dataset")\nplt.grid(
        alpha=0.3)\nplt.legend()\nplt.tight_layout()\n'
```

```
In [13]: def show_reconstructions(model, data_loader, device):
```

```

def show_reconstructions(model, data_loader, device):
    model.eval()
    with lucid.no_grad():
        for data, _ in data_loader:
            data = data.to(device)
            data = normalize(data)
            recon, _, _ = model(data)
            break


    imgs = data[:8].numpy()
    recons = recon[:8].numpy()

    import matplotlib.pyplot as plt
    _, axes = plt.subplots(2, 8, figsize=(16, 4))
    for i in range(8):
        axes[0, i].imshow(imgs[i].transpose((1, 2, 0)))
        axes[1, i].imshow(recons[i].transpose((1, 2, 0)))
        axes[0, i].axis('off')
        axes[1, i].axis('off')
    plt.suptitle("Originals (top) vs Reconstructions (bottom)")
    plt.show()

show_reconstructions(model, train_loader, device)

```

Originals (top) vs Reconstructions (bottom)



In [14]:

```

def generate_image_grid(latent_dim=128, grid_size=5):
    model.eval()
    with lucid.no_grad():
        # Sample z ~ N(0, 1)
        z = lucid.random.randn(grid_size * grid_size, latent_dim).to(device)
        out = model.decode(z).numpy()

        # Clamp to [0, 1] range in case of minor noise
        out = out.clip(0, 1)

        # Plot as 5x5 grid
        fig, axes = plt.subplots(grid_size, grid_size, figsize=(6, 6))
        for i in range(grid_size * grid_size):
            ax = axes[i // grid_size, i % grid_size]
            ax.imshow(out[i].transpose((1, 2, 0)))
            ax.axis('off')

        plt.tight_layout()
        plt.show()

generate_image_grid()

```

